

Perun: Keep Your Project's Performance Under Control

DevConf.cz Mini 2022

Tomas Fiedor, Jiri Pavela, Adam Rogalewicz, Tomas Vojnar

Brno University of Technology, Faculty of Information Technology

Supported by:

Red Hat, Inc.

Czech Science Foundation Project 20-07487S
JCMM PhD Talent Scholarship Programme



Motivation:

Why Care About Performance?

- Software performance bugs are an **omnipresent problem**¹:

¹Source: <https://accidentallyquadratic.tumblr.com/>

- Software performance bugs are an **omnipresent problem**¹:
 - **Cluster computing engine** may *freeze* after an update!



- An internal check for uniqueness
→ *hanging effectively forever* for large job batch.

¹Source: <https://accidentallyquadratic.tumblr.com/>

- Software performance bugs are an **omnipresent problem**¹:
 - **Cluster computing engine** may *freeze* after an update!
 - **Cloud services** may *crash*!



- An internal check for uniqueness
→ *hanging effectively forever* for large job batch.
- A regular expression for stripping whitespaces
→ *34 minutes long outage*.

¹Source: <https://accidentallyquadratic.tumblr.com/>

- Software performance bugs are an **omnipresent problem**¹:
 - **Cluster computing engine** may *freeze* after an update!
 - **Cloud services** may *crash*!
 - **Parsers** may experience significant *slowdown*!



- An internal check for uniqueness
→ *hanging effectively forever* for large job batch.



- A regular expression for stripping whitespaces
→ *34 minutes long outage*.



- One of *Chrome's* parsers
→ *noticeable slowdown* for long lines.

¹Source: <https://accidentallyquadratic.tumblr.com/>



- $\mathcal{O}(n^2)$ space implementation of C# constant folding
→ *compiler runs out of memory.*



- $\mathcal{O}(n^2)$ space implementation of C# constant folding
→ *compiler runs out of memory.*



- $\mathcal{O}(n^2)$ pattern matching algorithm in Elasticsearch
→ up to $\frac{1}{2}$ CPU-time spent in `Regex.simpleMatch`.



- $\mathcal{O}(n^2)$ space implementation of C# constant folding
→ *compiler runs out of memory.*



- $\mathcal{O}(n^2)$ pattern matching algorithm in Elasticsearch
→ up to $\frac{1}{2}$ CPU-time spent in `Regex.simpleMatch`.



- Array used for tags lookup in Vim
→ $\mathcal{O}(n^2)$ complexity in the number of matches.



- $\mathcal{O}(n^2)$ space implementation of C# constant folding
→ *compiler runs out of memory.*



- $\mathcal{O}(n^2)$ pattern matching algorithm in Elasticsearch
→ up to $\frac{1}{2}$ CPU-time spent in `Regex.simpleMatch`.



- Array used for tags lookup in Vim
→ $\mathcal{O}(n^2)$ complexity in the number of matches.



- Godoc source code parsing
→ $\mathcal{O}(n^2)$ loop for Go structs definitions.

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average less time to fix;
 - can be fixed by less experienced developers;
 - the fix is generally smaller.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. 8 days);
 - can be fixed by less experienced developers;
 - the fix is generally smaller.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. 8 days);
 - can be fixed by **less experienced developers**;
 - the fix is generally smaller.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. **8** days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. **19** LoC).

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. **8** days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. **19** LoC).
- Hence new bugs should be discovered **as soon as possible**.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. **8** days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. **19** LoC).
 - Hence new bugs should be discovered **as soon as possible**.
- ⇒ **Continuous Integration (CI) and Automated Testing!**

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. **8** days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. **19** LoC).
 - Hence new bugs should be discovered **as soon as possible**.
- ⇒ **Continuous Integration (CI) and Automated Testing!**
- Already commonly utilized for testing the project's functionality.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. **8** days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. **19** LoC).
 - Hence new bugs should be discovered **as soon as possible**.
- ⇒ **Continuous Integration (CI) and Automated Testing!**
- Already commonly utilized for testing the project's **functionality**.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. **8** days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. **19** LoC).
 - Hence new bugs should be discovered **as soon as possible**.
- ⇒ **Continuous Integration (CI) and Automated Testing!**
- Already commonly utilized for testing the project's **functionality**.
 - However, what about **performance** bugs?
 - Most solutions are either ad-hoc or proprietary.

²T.-H. Chen et al.: *An empirical study of dormant bugs*

- **Recency is important as well:** it pays off to discover bugs quickly.
 - Recently introduced bugs, as opposed to dormant bugs²,
 - take on average **less time to fix** (5 vs. 8 days);
 - can be fixed by **less experienced developers**;
 - the fix is generally **smaller** (10 vs. 19 LoC).
 - Hence new bugs should be discovered **as soon as possible**.
- ⇒ **Continuous Integration (CI) and Automated Testing!**
- Already commonly utilized for testing the project's **functionality**.
 - However, what about **performance** bugs?
 - Most solutions are either ad-hoc or proprietary.
 - We are not aware of any **complex open-source** solutions.

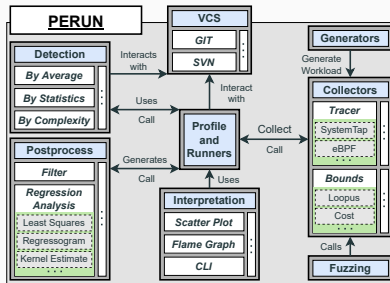
²T.-H. Chen et al.: *An empirical study of dormant bugs*

Meet Perun: **Performance Version System**

Perun³ = Complex Solution for Performance Analysis and Testing

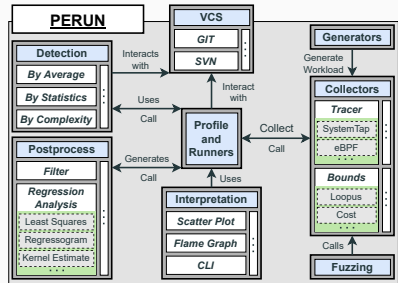
³Available at: <https://github.com/tfiedor/perun/>

Perun³ = Complex Solution for Performance Analysis and Testing =
 = **Collects** performance data



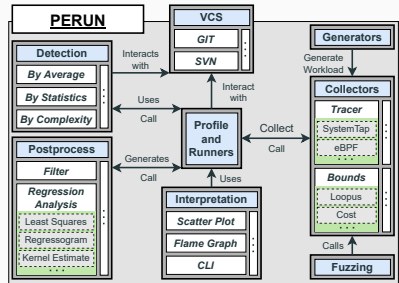
³Available at: <https://github.com/tfiedor/perun/>

- Perun**³ = Complex Solution for Performance Analysis and Testing =
- = **Collects** performance data
 - + **Creates** performance models
 - Constant c , linear $an + b$, ...



³Available at: <https://github.com/tfiedor/perun/>

- Perun**³ = Complex Solution for Performance Analysis and Testing =
- = **Collects** performance data
 - + **Creates** performance models
 - Constant c , linear $an + b$, ...
 - + **Integrates** VCS
 - Access to project history.



³Available at: <https://github.com/tfiedor/perun/>

Perun³ = Complex Solution for Performance Analysis and Testing =

= **Collects** performance data

+ **Creates** performance models

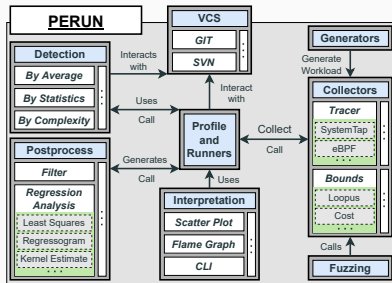
- Constant c , linear $an + b$, ...

+ **Integrates** VCS

- Access to project history.

+ **Detects** performance changes

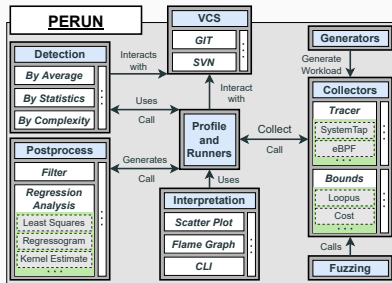
- Degradations, optimizations.



³Available at: <https://github.com/tfiedor/perun/>

Perun³ = Complex Solution for Performance Analysis and Testing =

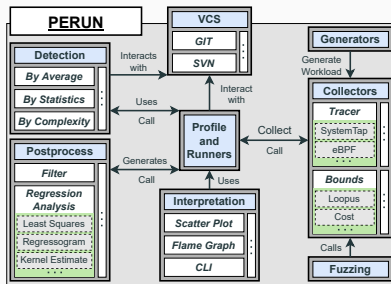
- = **Collects** performance data
- + **Creates** performance models
 - Constant c , linear $an + b$, ...
- + **Integrates** VCS
 - Access to project history.
- + **Detects** performance changes
 - Degradations, optimizations.
- + **Visualizes** performance



³Available at: <https://github.com/tfiedor/perun/>


Perun³ = Complex Solution for Performance Analysis and Testing =

- = **Collects** performance data*
- + **Creates** performance models
 - Constant c , linear $an + b$, ...
- + **Integrates** VCS
 - Access to project history.
- + **Detects** performance changes
 - Degradations, optimizations.
- + **Visualizes** performance




* Often the only thing done by traditional profilers.


³Available at: <https://github.com/tfiedor/perun/>




- **Four** major steps: Repository → Profiles → Models → Detection




- **Four** major steps: **Repository** → Profiles → Models → Detection



- **Four** major steps: Repository → Profiles → Models → Detection



- **Four** major steps: Repository → Profiles → Models → Detection



- **Four** major steps: Repository → Profiles → Models → Detection

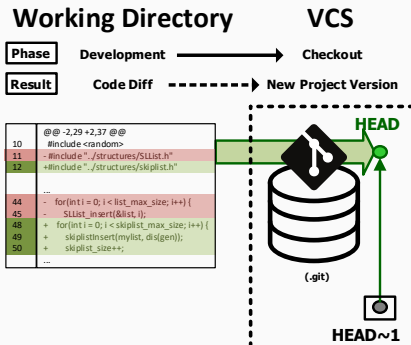
Working Directory

Phase Development

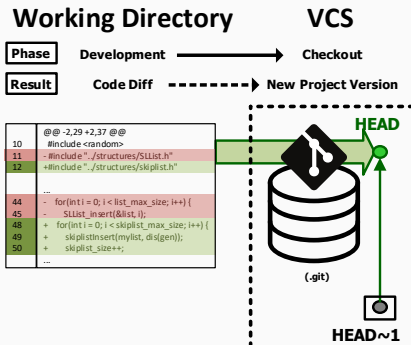
Result Code Diff

	@@ -2,29 +2,37 @@
10	#include <random>
11	- #include "../structures/SList.h"
12	+ #include "../structures/skiplist.h"
	...
44	- for(int i = 0; i < list_max_size; i++) {
45	- SList_insert(&list, i);
48	+ for(int i = 0; i < skiplist_max_size; i++) {
49	+ skiplistinsert(mylist, dis(gen));
50	+ skiplist_size++;
	...

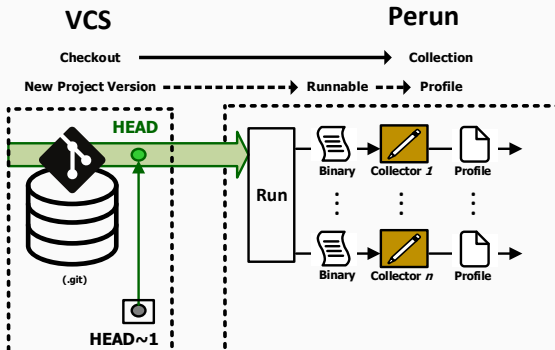
1. We create the project's **working directory**.



1. We create the project's **working directory**.
2. We initialize a **VCS** (e.g., Git) for project versioning.




1. We create the project's **working directory**.
2. We initialize a **VCS** (e.g., Git) for project versioning.
3. We initialize **Perun** in the repository alongside the VCS.



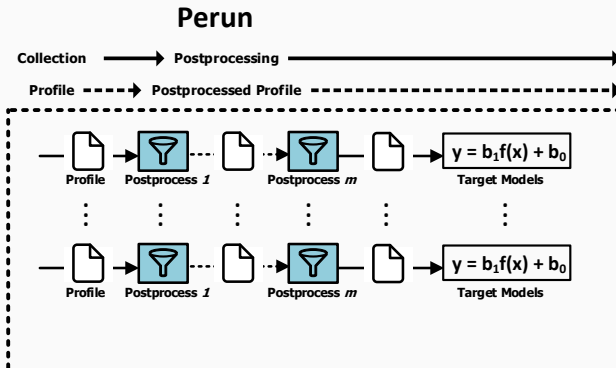
4. We measure project's performance and obtain **profiles**.

- Profiles are stored within Perun and linked to the corresponding VCS version (e.g., commit).




4. We measure project's performance and obtain **profiles**.

- Profiles are stored within Perun and **linked to the corresponding VCS version** (e.g., commit).




5. We create **performance models** from profiles using postprocessors.
- Models are stored within Perun alongside the profiles.



5. We create **performance models** from profiles using postprocessors.
- Models are **stored within Perun** alongside the profiles.


- Models in Perun are **mathematical functions** of the input size or **statistical summaries** describing the main features of the profile.

Plot of 'amount' per 'structure-unit-size'; uid: SLList_search(SLList*, int); method: full; interval <0, 11892




- Models in Perun are **mathematical functions** of the input size or **statistical summaries** describing the main features of the profile.
- For example, in **regression analysis**:

Plot of 'amount' per 'structure-unit-size'; uid: SLList_search(SLList*, int); method: full; interval <0, 11892




- Models in Perun are **mathematical functions** of the input size or **statistical summaries** describing the main features of the profile.
- For example, in **regression analysis**:
 - the curve is described using **function** $y = b_1 f(x) + b_0$,

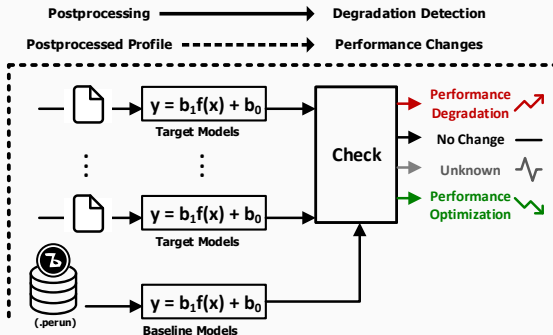
Plot of 'amount' per 'structure-unit-size'; uid: SLList_search(SLList*, int); method: full; interval <0, 11892



- Models in Perun are **mathematical functions** of the input size or **statistical summaries** describing the main features of the profile.
- For example, in **regression analysis**:
 - the curve is described using **function** $y = b_1 f(x) + b_0$,
 - and the model quality as **coefficient of determination** R^2 .

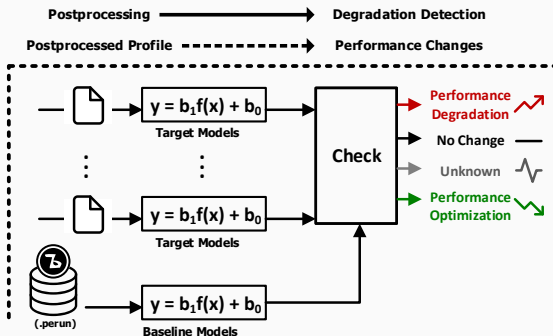
of 'amount' per 'structure-unit-size'; uid: SLList_search(SLList*, int); method: initial_guess; interval <0, '





6. We detect **performance changes** using models or directly profiles.

- Target refers to the current version.
- Baseline refers to the previous version used for comparison.



6. We detect **performance changes** using models or directly profiles.

- **Target** refers to the current version.
- **Baseline** refers to the previous version used for comparison.

- Multiple **detection algorithms** are implemented in Perun:
 - *Best Model Order Equality*
 - *Integral Comparison*
 - ...
 - *Exclusive-Time Outliers*

- Multiple **detection algorithms** are implemented in Perun:
 - *Best Model Order Equality*
 - *Integral Comparison*
 - ...
 - *Exclusive-Time Outliers*

- Multiple **detection algorithms** are implemented in Perun:
 - *Best Model Order Equality*
 - *Integral Comparison*
 - ...
 - *Exclusive-Time Outliers*
- **Exclusive-Time Outliers:**
 - Based on per-function comparison of exclusive-time **deltas** Δ .
 - *Exclusive-time (self)*: time spent exclusively in a function.

- Multiple **detection algorithms** are implemented in Perun:
 - *Best Model Order Equality*
 - *Integral Comparison*
 - ...
 - *Exclusive-Time Outliers*
- **Exclusive-Time Outliers:**
 - Based on per-function comparison of exclusive-time **deltas** Δ .
 - *Exclusive-time (self)*: time spent exclusively in a function.
 - A hierarchy of outlier detection methods determines the **severity**:
 - **Modified Z-score**

$$y_i = \frac{x_i - \tilde{X}}{k \cdot \text{median}(|x_j - \tilde{X}|)}$$

- Multiple **detection algorithms** are implemented in Perun:
 - *Best Model Order Equality*
 - *Integral Comparison*
 - ...
 - *Exclusive-Time Outliers*
- **Exclusive-Time Outliers:**
 - Based on per-function comparison of exclusive-time **deltas** Δ .
 - *Exclusive-time (self)*: time spent exclusively in a function.
 - A hierarchy of outlier detection methods determines the **severity**:
 - **Modified Z-score**
 - **IQR multiple**

$$y_i = \frac{x_i - \tilde{X}}{k \cdot \text{median}(|x_j - \tilde{X}|)}$$
$$Q_1 - k \cdot IQR < x < Q_3 + k \cdot IQR$$

- Multiple **detection algorithms** are implemented in Perun:
 - *Best Model Order Equality*
 - *Integral Comparison*
 - ...
 - *Exclusive-Time Outliers*
- **Exclusive-Time Outliers:**
 - Based on per-function comparison of exclusive-time **deltas** Δ .
 - *Exclusive-time (self)*: time spent exclusively in a function.
 - A hierarchy of outlier detection methods determines the **severity**:
 - **Modified Z-score**
 - **IQR multiple**
 - **Standard deviation multiple**

$$y_i = \frac{x_i - \tilde{X}}{k \cdot \text{median}(|x_j - \tilde{X}|)}$$
$$Q_1 - k \cdot IQR < x < Q_3 + k \cdot IQR$$
$$\bar{X} - k \cdot \sigma < x < \bar{X} + k \cdot \sigma$$

Demonstration of Perun #1: Finding Performance Changes

- **CPython:** Reference C implementation of a Python interpreter.

⁴Reported by user mdboom: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356**⁴: A performance **regression** in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).

⁴Reported by user mdboom: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356**⁴: A performance **regression** in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).
 - Replicated using the pypperformance ctypes benchmark.

⁴Reported by user mdboom: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356**⁴: A performance **regression** in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).
 - Replicated using the pyperformance ctypes benchmark.
 - **Fixed soon after the report.**

⁴Reported by user mdboom: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356⁴**: A performance regression in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).
 - Replicated using the `pyperformance` `ctypes` benchmark.
 - **Fixed soon after the report.**

So problem solved. Why use Perun?

⁴Reported by user `mdboom`: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356⁴**: A performance regression in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).
 - Replicated using the `pyperformance` `ctypes` benchmark.
 - Fixed soon after the report.

So problem solved. Why use Perun?

- Discovering the issue and finding the root cause is the hard part.

⁴Reported by user `mdboom`: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356**⁴: A performance regression in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).
 - Replicated using the `pyperformance` `ctypes` benchmark.
 - Fixed soon after the report.

So problem solved. Why use Perun?

- **Discovering the issue** and **finding the root cause** is the hard part.
 - Often requires **significant** manual effort by the developers.

⁴Reported by user mdboom: <https://github.com/python/cpython/issues/92356>

- **CPython**: Reference C implementation of a Python interpreter.
- **Issue #92356**⁴: A performance regression in ctypes module.
 - $\approx 8\%$ higher function call overhead (py3.11.0a7 vs. py3.10.4).
 - Replicated using the `pyperformance` `ctypes` benchmark.
 - Fixed soon after the report.

So problem solved. Why use Perun?

- **Discovering the issue** and **finding the root cause** is the hard part.
 - Often requires **significant** manual effort by the developers.
- **Perun** reduces this effort and helps the developers.
 - Perun utilizes the **recency** principle and results from **past profiling**.

⁴Reported by user `mdboom`: <https://github.com/python/cpython/issues/92356>

Using **Perun**, we could handle the **issue #92356** as follows:

Using **Perun**, we could handle the **issue #92356** as follows:

1. We **initialize** a CPython repository with Perun.

Perun commands

```
perun init
```

Using **Perun**, we could handle the **issue #92356** as follows:

1. We **initialize** a CPython repository with Perun.
2. We **store** a profile for CPython **v3.10.4** ctypes benchmark in Perun.

Perun commands

```
perun init
```


Using **Perun**, we could handle the **issue #92356** as follows:

1. We **initialize** a CPython repository with Perun.
2. We **store** a profile for CPython **v3.10.4** ctypes benchmark in Perun.
 - We denote this profile as **baseline**.
 - Perun handles the *profile-commit* link internally.

Perun commands

```
perun init
```

3. CPython v3.11.0a7 **rolls out**.

3. CPython v3.11.0a7 **rolls out**.
4. We **profile** the ctypes benchmark for CPython **v3.11.0a7**.

Perun commands

```
perun collect -c <py3.11.0a7> -a <benchmark> trace -b <files>
```

3. CPython v3.11.0a7 **rolls out**.
4. We **profile** the ctypes benchmark for CPython **v3.11.0a7**.
 - We denote the resulting profile as **target**.

Perun commands

```
perun collect -c <py3.11.0a7> -a <benchmark> trace -b <files>  
perun add <target>
```

3. CPython v3.11.0a7 **rolls out**.
4. We **profile** the ctypes benchmark for CPython **v3.11.0a7**.
 - We denote the resulting profile as **target**.
5. We **compare** the **baseline** and **target** profiles.

Perun commands

```
perun collect -c <py3.11.0a7> -a <benchmark> trace -b <files>  
perun add <target>  
perun check -f profiles <baseline> <target>
```

3. CPython v3.11.0a7 **rolls out**.
4. We **profile** the ctypes benchmark for CPython **v3.11.0a7**.
 - We denote the resulting profile as **target**.
5. We **compare** the **baseline** and **target** profiles.
 - Perun supports multiple comparison algorithms.
 - For this particular issue, we used *Exclusive-Time Outliers*.

Perun commands

```
perun collect -c <py3.11.0a7> -a <benchmark> trace -b <files>
perun add <target>
perun check -f profiles <baseline> <target>
```

Location		Result	TΔ [ms]	TΔ [%]
<code>_ctypes_init_fielddesc</code>		NotInBaseline	77.95	5.23
<code>_ctypes_get_fielddesc</code>		SevereDegradation	52.9	3.55
<code>_ctypes_callproc</code>		Degradation	2.84	0.19
		...		
<code>_ctypes.cpython-311</code>		TotalDegradation	136.92	9.19

* TΔ: exclusive-time delta of *target* – *baseline*.

Location	Result	TΔ [ms]	TΔ [%]
<code>_ctypes_init_fielddesc</code>	NotInBaseline	77.95	5.23
<code>_ctypes_get_fielddesc</code>	SevereDegradation	52.9	3.55
<code>_ctypes_callproc</code>	Degradation	2.84	0.19
...			
<code>_ctypes.cpython-311</code>	TotalDegradation	136.92	9.19

* TΔ: exclusive-time delta of *target* – *baseline*.

Location	Result	TΔ [ms]	TΔ [%]
<code>_ctypes_init_fielddesc</code>	NotInBaseline	77.95	5.23
<code>_ctypes_get_fielddesc</code>	SevereDegradation	52.9	3.55
<code>_ctypes_callproc</code>	Degradation	2.84	0.19
...			
<code>_ctypes.cpython-311</code>	TotalDegradation	136.92	9.19

* TΔ: exclusive-time delta of *target* – *baseline*.

- **Root cause** of the issue: **repeated calls** of an init function.

Function `_ctypes_get_fielddesc`

```
if (!initialized) {  
    _ctypes_init_fielddesc();  
}
```

6. We **create** a new hotfix branch and **fix** the issue.

Fixing `_ctypes_get_fielddesc`

```
if (!initialized) {  
+     initialized = 1;  
    _ctypes_init_fielddesc();  
}
```

6. We **create** a new hotfix branch and **fix** the issue.
7. We **Profile** the CPython hotfixed version.

Fixing `_ctypes_get_fielddesc`

```
if (!initialized) {  
+     initialized = 1;  
    _ctypes_init_fielddesc();  
}
```

Perun commands

```
perun collect -c <py3.11.0a7-fix> -a <benchmark> trace <...>
```

6. We **create** a new hotfix branch and **fix** the issue.
7. We **Profile** the CPython hotfixed version.
 - We denote the resulting profile as **hotfix**.

Fixing `_ctypes_get_fielddesc`

```
if (!initialized) {  
+   initialized = 1;  
    _ctypes_init_fielddesc();  
}
```

Perun commands

```
perun collect -c <py3.11.0a7-fix> -a <benchmark> trace <...>  
perun add <hotfix>
```

6. We **create** a new hotfix branch and **fix** the issue.
7. We **Profile** the CPython hotfixed version.
 - We denote the resulting profile as **hotfix**.
8. We **compare** the **baseline** and **hotfix** profiles.

Fixing `_ctypes_get_fielddesc`

```
if (!initialized) {  
+   initialized = 1;  
    _ctypes_init_fielddesc();  
}
```

Perun commands

```
perun collect -c <py3.11.0a7-fix> -a <benchmark> trace <...>  
perun add <hotfix>  
perun check -f profiles <baseline> <hotfix>
```

Location	Result	Δ [ms]	Δ_{old} [ms]	Δ [%]	Δ_{old} [%]
...					
<code>_ctypes_get_fielddesc</code>	MaybeDegradation	0.89	52.9	0.06	3.55
<code>_ctypes_init_fielddesc</code>	NotInBaseline	0.02	77.95	0.00	5.23
<code>_ctypes.cpython-311</code>	TotalDegradation	23.45	136.92	1.70	9.19

* Δ : exclusive-time delta of *hotfix*—*baseline*.

* Δ_{old} : exclusive-time delta of *target*—*baseline*.

Location		Result	Δ [ms]	Δ_{old} [ms]	Δ [%]	Δ_{old} [%]
...						
<code>_ctypes_get_fielddesc</code>	MaybeDegradation		0.89	52.9	0.06	3.55
<code>_ctypes_init_fielddesc</code>	NotInBaseline		0.02	77.95	0.00	5.23
<code>_ctypes.cpython-311</code>	TotalDegradation		23.45	136.92	1.70	9.19

* Δ : exclusive-time delta of *hotfix*—*baseline*.

* Δ_{old} : exclusive-time delta of *target*—*baseline*.

- The `_ctypes_get_fielddesc` Δ has improved significantly.

Location		Result	Δ [ms]	Δ_{old} [ms]	Δ [%]	Δ_{old} [%]
...						
<code>_ctypes_get_fielddesc</code>	MaybeDegradation		0.89	52.9	0.06	3.55
<code>_ctypes_init_fielddesc</code>	NotInBaseline		0.02	77.95	0.00	5.23
<code>_ctypes.cpython-311</code>	TotalDegradation		23.45	136.92	1.70	9.19

* Δ : exclusive-time delta of *hotfix*—*baseline*.

* Δ_{old} : exclusive-time delta of *target*—*baseline*.

- The `_ctypes_get_fielddesc` Δ has improved significantly.
- The `_ctypes_init_fielddesc` Δ is now negligible.

Location	Result	Δ [ms]	Δ_{old} [ms]	Δ [%]	Δ_{old} [%]
...					
<code>_ctypes_get_fielddesc</code>	MaybeDegradation	0.89	52.9	0.06	3.55
<code>_ctypes_init_fielddesc</code>	NotInBaseline	0.02	77.95	0.00	5.23
<code>_ctypes.cpython-311</code>	TotalDegradation	23.45	136.92	1.70	9.19

* Δ : exclusive-time delta of *hotfix*—*baseline*.

* Δ_{old} : exclusive-time delta of *target*—*baseline*.

- The `_ctypes_get_fielddesc` Δ has improved significantly.
- The `_ctypes_init_fielddesc` Δ is now negligible.

⇒ Perun leverages **VCS** and **Recency** to successfully discover and help locate performance issues in new project versions.

Demonstration of Perun #2: Generating Workloads

- Recall the **Stack Overflow** issue:



- A regular expression for stripping whitespaces
→ *34 minutes long outage.*

- Recall the **Stack Overflow** issue:



stack**overflow**

- A regular expression for stripping whitespaces
→ *34 minutes long outage.*

The offending regular expression:

```
^[\s\u200c]+|[\s\u200c]+$
```

- Recall the **Stack Overflow** issue:

stack**overflow**

- A regular expression for stripping whitespaces
→ *34 minutes long outage.*

The offending regular expression:

```
^[\s\u200c]+|[\s\u200c]+$
```

- This Regex can lead to **extensive backtracking**.
- Can **Perun** help us with detecting this **potential performance issue**?

- Recall the **Stack Overflow** issue:

stack**overflow**

- A regular expression for stripping whitespaces
→ *34 minutes long outage.*

The offending regular expression:

```
^[\s\u200c]+|[\s\u200c]+$
```

- This Regex can lead to **extensive backtracking**.
- Can **Perun** help us with detecting this **potential performance issue**?

Perun's Performance Fuzzing!⁵

⁵Builds on a principle originally proposed by C. Lemieux et al.:

PerfFuzz: automatically generating pathological inputs.

- **Fuzzing (Fuzz testing)** is a form of **fault injection stress testing**.

- **Fuzzing (Fuzz testing)** is a form of **fault injection stress testing**.
 - Malformed inputs are generated and fed to the program.
 - The program response is then monitored for **crashes** or **errors**.

- **Fuzzing (Fuzz testing)** is a form of **fault injection stress testing**.
 - Malformed inputs are generated and fed to the program.
 - The program response is then monitored for **crashes** or **errors**.
- **Perun's Fuzzer⁶** is **profiling-guided**.

⁶M. Liscinsky: *Fuzz Testing of Program Performance*.

- **Fuzzing (Fuzz testing)** is a form of **fault injection stress testing**.
 - Malformed inputs are generated and fed to the program.
 - The program response is then monitored for **crashes** or **errors**.
- **Perun's Fuzzer⁶** is **profiling-guided**.
 - The program is not monitored for errors or crashes, but **slowdown**.
 - We evaluate the results using profiling.

⁶M. Liscinsky: *Fuzz Testing of Program Performance*.

- **Fuzzing (Fuzz testing)** is a form of **fault injection stress testing**.
 - Malformed inputs are generated and fed to the program.
 - The program response is then monitored for **crashes** or **errors**.
- **Perun's Fuzzer**⁶ is **profiling-guided**.
 - The program is not monitored for errors or crashes, but **slowdown**.
 - We evaluate the results using profiling.
 - **The goal** is to find inputs that cause **severe slowdown**⁷.

⁶M. Liscinsky: *Fuzz Testing of Program Performance*.

⁷In some of our experiments, Perun's Fuzzer achieved a **slowdown of several hours!**

- For this experiment, we used the following **settings**:

- For this experiment, we used the following **settings**:
 - **Profiled program**: 25 LoC C++ program matching the offending regex (using the `regex_search` function) in the input file.

- For this experiment, we used the following **settings**:
 - **Profiled program**: 25 LoC C++ program matching the offending regex (using the `regex_search` function) in the input file.
 - **Seed**: 150 LoC C implementation of a parallel grep.

- For this experiment, we used the following **settings**:
 - **Profiled program**: 25 LoC C++ program matching the offending regex (using the `regex_search` function) in the input file.
 - **Seed**: 150 LoC C implementation of a parallel grep.
 - **Perun's own mutation rules**. Particularly useful rule is:
 - T.10: Insert whitespaces to a random position in a string.

- For this experiment, we used the following **settings**:
 - **Profiled program**: 25 LoC C++ program matching the offending regex (using the `regex_search` function) in the input file.
 - **Seed**: 150 LoC C implementation of a parallel grep.
 - **Perun's own mutation rules**. Particularly useful rule is:
 - T.10: Insert whitespaces to a random position in a string.
 - **Size limits**: (a) 5 000 bytes, (b) 10 000 bytes.

Input	Size [B]	Lines	Whitespaces	Duration [s]	Slowdown
<i>seed</i>	3535	150	306	0.096	—
<i>worst-case_a</i>	5000	5	4881	1.566	16.3x
<i>worst-case_b</i>	10000	17	9603	2.611	27.2x

* We let the fuzzing run for several hours to obtain the shown workloads.

* The generated workloads can then be used for future performance testing.

Input	Size [B]	Lines	Whitespaces	Duration [s]	Slowdown
<i>seed</i>	3535	150	306	0.096	—
<i>worst-case_a</i>	5000	5	4881	1.566	16.3x
<i>worst-case_b</i>	10000	17	9603	2.611	27.2x

* We let the fuzzing run for several hours to obtain the shown workloads.

* The generated workloads can then be used for future performance testing.

Input	Size [B]	Lines	Whitespaces	Duration [s]	Slowdown
<i>seed</i>	3535	150	306	0.096	—
<i>worst-case_a</i>	5000	5	4881	1.566	16.3x
<i>worst-case_b</i>	10000	17	9603	2.611	27.2x

* We let the fuzzing run for several hours to obtain the shown workloads.

* The generated workloads can then be used for future performance testing.

Input	Size [B]	Lines	Whitespaces	Duration [s]	Slowdown
<i>seed</i>	3535	150	306	0.096	—
<i>worst-case_a</i>	5000	5	4881	1.566	16.3x
<i>worst-case_b</i>	10000	17	9603	2.611	27.2x

* We let the fuzzing run for several hours to obtain the shown workloads.

* The generated workloads can then be used for future performance testing.

⇒ Perun Fuzzer can force potential performance issues to manifest!

Input	Size [B]	Lines	Whitespaces	Duration [s]	Slowdown
<i>seed</i>	3535	150	306	0.096	—
<i>worst-case_a</i>	5000	5	4881	1.566	16.3x
<i>worst-case_b</i>	10000	17	9603	2.611	27.2x

* We let the fuzzing run for several hours to obtain the shown workloads.


* The generated workloads can then be used for future performance testing.

⇒ Perun Fuzzer can force potential performance issues to manifest!

- We employ different **mutation strategies** based on the input type.
 - Text files, binary files, domain-specific (e.g., XML), ...


Ongoing and Future Work

- We focus on increasing profiling **granularity**.
 - Measuring time spent per function **basic block**⁸.




⁸P. Mocary: *Performance Analysis of Programs Based on PIN Framework*.

- Moreover, we focus on increasing profiling **precision**.
- Measuring time spent w.r.t. function **parameter values**⁹.




⁹P. Mocary: *Performance Analysis of Programs Based on PIN Framework*.

- Finally, we focus on increasing profiling **efficiency**.
 - Utilizing a collection of heuristics to **speed up** the profiling¹⁰.
 - We propose to select and sample particularly **important** functions.



¹⁰J. Pavela: *Efficient Techniques for Program Performance Analysis*.


- Finally, we focus on increasing profiling **efficiency**.
 - Utilizing a collection of heuristics to **speed up** the profiling¹⁰.
 - We propose to select and sample particularly **important** functions.



- The main **challenge**: achieving **sufficient precision**.

¹⁰J. Pavela: *Efficient Techniques for Program Performance Analysis*.

- Finally, we focus on increasing profiling **efficiency**.
 - Utilizing a collection of heuristics to **speed up** the profiling¹⁰.
 - We propose to select and sample particularly **important** functions.



- The main **challenge**: achieving **sufficient precision**.
 - ⇒ Fully profiling all the important functions.

¹⁰J. Pavela: *Efficient Techniques for Program Performance Analysis*.

- **Future work:**

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...
 - More performance **metrics**:
 - Currently mainly time consumption of functions.
 - Memory consumption in experimental stage.

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...
 - More performance **metrics**:
 - Currently mainly time consumption of functions.
 - Memory consumption in experimental stage.
 - ⇒ Energy consumption, cache hit/miss, scheduler, ...

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...
 - More performance **metrics**:
 - Currently mainly time consumption of functions.
 - Memory consumption in experimental stage.
 - ⇒ Energy consumption, cache hit/miss, scheduler, ...

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...
 - More performance **metrics**:
 - Currently mainly time consumption of functions.
 - Memory consumption in experimental stage.
 - ⇒ Energy consumption, cache hit/miss, scheduler, ...
 - Supporting well-established and/or state-of-the-art **tools**:
 - Currently Facebook Infer Cost, Loopus.

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...
 - More performance **metrics**:
 - Currently mainly time consumption of functions.
 - Memory consumption in experimental stage.
 - ⇒ Energy consumption, cache hit/miss, scheduler, ...
 - Supporting well-established and/or state-of-the-art **tools**:
 - Currently Facebook Infer Cost, Loopus.
 - ⇒ Valgrind tool suite, Perf, Gprof, ...

- **Future work:**
 - Support for **more languages**:
 - Currently mainly C/C++.
 - ⇒ C#, Java, Python, ...
 - More performance **metrics**:
 - Currently mainly time consumption of functions.
 - Memory consumption in experimental stage.
 - ⇒ Energy consumption, cache hit/miss, scheduler, ...
 - Supporting well-established and/or state-of-the-art **tools**:
 - Currently Facebook Infer Cost, Loopus.
 - ⇒ Valgrind tool suite, Perf, Gprof, ...
 - Performance analysis of **Dynamic Data Structures**

Conclusion

- **Perun** = Complex Performance Analysis and Testing Solution.

- **Perun** = Complex Performance Analysis and Testing Solution.
 - Integrates VCS.
 - Collects performance data.
 - Derives performance models.
 - Detects performance changes.
 - Visualises performance.

- **Perun** = Complex Performance Analysis and Testing Solution.
 - Integrates VCS.
 - Collects performance data.
 - Derives performance models.
 - Detects performance changes.
 - Visualises performance.
- ⇒ Not just mere profiling!

- **Perun** = Complex Performance Analysis and Testing Solution.
 - Integrates VCS.
 - Collects performance data.
 - Derives performance models.
 - Detects performance changes.
 - Visualises performance.

⇒ Not just mere profiling!
- Nicely complements cheap benchmarking.
 - Shown on a real-world CPython performance bug.
- Forces potential performance issues to manifest.
 - Replication of a Stack Overflow regex performance bug.

- **Perun** = **Complex Performance Analysis and Testing Solution**.

- **Integrates** VCS.
- **Collects** performance data.
- **Derives** performance models.
- **Detects** performance changes.
- **Visualises** performance.

⇒ **Not just mere profiling!**

- **Nicely complements cheap benchmarking.**

- Shown on a **real-world** CPython performance bug.

- **Forces potential performance issues to manifest.**

- Replication of a Stack Overflow **regex performance bug**.

- **Ongoing and Future work:**

- Improving **granularity**, **precision** and **efficiency** of profiling.
- Support for more **languages**, performance **metrics**, existing **tools**.

Perun: Keep Your Project's Performance Under Control

DevConf.cz Mini 2022

Tomas Fiedor, Jiri Pavela, Adam Rogalewicz, Tomas Vojnar

Brno University of Technology, Faculty of Information Technology

Supported by:

Red Hat, Inc.

Czech Science Foundation Project 20-07487S
JCMM PhD Talent Scholarship Programme

